

# **LINEAR REGRESSION IV**

11.11.2020

# RECAP

- \* receptive fields
- \* what kind of stuff in the world does a neuron (or voxel, or whatever) respond to?
- \* we can estimate receptive fields by measuring responses in different conditions, and then fitting a model using **regression**

# RECAP

- \* we need to find weights that minimize squared error

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$

- \* one method is gradient descent
- \* another is the analytic solution  
(`np.linalg.lstsq`)

# RECAP

- \* how do we measure *how good* a regression model is?
- \*  $R^2$  aka the **coefficient of determination** or **variance explained**
- \*  $R^2 = 1 - (\text{RSS} / \text{TSS})$ 
  - \* RSS is “residual sum of squares” (=squared error)
  - \* TSS is “total sum of squares” (error with  $\beta=0$ )

# EVALUATING REGRESSION MODELS


- \* suppose that we are given a matrix of variables (aka regressors)  $X$ , and a vector of outputs  $Y$
- \* we fit a linear model  $\hat{Y} = X \cdot \beta$
- \* then we evaluate it by computing  $R^2$  using  $X$  and  $Y$
- \* what are the possible values of  $R^2$ ?

# IN-SET VS. OUT-OF-SET EVALUATION

- \* evaluating a regression model using the same data that we used to train/estimate/fit it is called *in-set evaluation*
- \* in-set evaluation is biased *upward*, and the amount of bias depends on the number of regressors in the model

# IN-SET VS. OUT-OF-SET EVALUATION

- \* for example: suppose we have  $N$  data points and  $N$  regressors that are pure noise—they have no relationship to the output whatsoever
- \* in-set variance explained is ***EXACTLY 1.0***
- \* ***THE MODEL IS PERFECT***



# IN-SET VS. OUT-OF-SET EVALUATION

- \* instead, what if you split up your  $X$  and  $Y$  into “training” and “test” sets?
- \* you could fit your regression model using  $(X_{\text{trn}}, Y_{\text{trn}})$ , and then test how well it works on  $(X_{\text{test}}, Y_{\text{test}})$ !
- \* is  $R^2$  biased in this case? What possible values can it take?



# REGRESSION STABILITY



- \* so ordinary least squares regression has a problem:
- \* if two regressors are *very similar* (i.e. correlated), then there are many weight combinations that would give ~the same answer!
- \* which of these combinations is “best” ends up being totally determined by *noise* (example)

# REGRESSION STABILITY


- \* this is bad: if your weights are essentially random, they are ~impossible to interpret, and model performance can suffer, so:
- \* (1) let's figure out when this is happening, and
- \* (2) let's stop it from happening

# REGRESSION STABILITY

- \* how do we know when regression is unstable?
- \* it's related to the shape of the error function!




# REGRESSION STABILITY







# REGRESSION STABILITY





# REGRESSION STABILITY



# REGRESSION STABILITY



# REGRESSION STABILITY





# REGRESSION STABILITY

- \* the dimensions of the error ellipse,  $k_1$  and  $k_2$ , are related to the **singular values** returned by `np.linalg.lstsq`!
- \* if the singular values (ordered from largest to smallest) are  $s_1, s_2$ , etc.,
- \* then  $k_1 \propto s_1^{-1}$ ,  $k_2 \propto s_2^{-1}$ , etc.

# REGRESSION STABILITY

- \* so it's easy to detect when a regression is unstable: look for tiny singular values! (example)
- \* (or at least, tiny relative to the largest singular value)

# REGRESSION STABILITY

- \* but what do you do if the regression is unstable?
- \* how could you possibly solve this problem?

**END**