

# LINEAR REGRESSION VI

11.16.2020

# **HOMEWORK 5**

\* due FRIDAY

# LINEAR REGRESSION

## READINGS

- \* Chapter 5 from PDSH, particularly 5.6 - “In Depth: Linear Regression”
- \* Chapter 16 from Inferential Thinking - “Inference for Regression”

# RECAP

- \* `np.linalg.lstsq` – numpy function that does least squares regression (often bad)
- \*  $R^2$  is a measure of how good a regression model is
- \* in-set vs. out-of-set evaluation of a regression model

# REGULARIZATION

- \* we modify the error function to be the sum of a **loss term** and a **penalty term**

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

# REGULARIZATION

- \* it also introduces an extra parameter,  $\lambda$ , which is the *regularization coefficient*, or, in this case, *ridge coefficient*

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

# RIDGE REGRESSION

- \* this type of regularization (penalizing the sum of squared weights) is called **ridge regression**
- \* and because the ridge error function (loss + penalty) is parabolic, it has an analytic solution!
- \* a nice implementation is in scikit-learn as **`sklearn.linear_model.Ridge`**

# RIDGE REGRESSION

- \* but when doing ridge regression you have a new issue: how do you choose the ridge parameter,  $\lambda$ ?
- \* if you train and test your regression model on the same piece of data,  $\lambda=0$  is always going to be the best
- \* *~bogus~*



# RIDGE REGRESSION

- \* if you train and test on different datasets (as discussed earlier) it's better
- \* but using your test data multiple times (to choose a parameter!) creates an issue of bias (aka **overfitting**)

# RIDGE REGRESSION

- \* the correct solution is **cross-validation**:
- \* break your dataset into training and test ( $X \rightarrow [X_{\text{trn}}, X_{\text{test}}]$ )
- \* further break up your training set ( $X_{\text{trn}} \rightarrow [X_{\text{fit}}, X_{\text{val}}]$ )
- \* fit weights using  $X_{\text{fit}}$ , choose  $\lambda$  based on performance on  $X_{\text{val}}$ , then finally test on  $X_{\text{test}}$

# RIDGE REGRESSION

- \* In reality you should understand this process (& its philosophical underpinnings), but you probably won't need to implement it yourself
- \* sklearn provides functions that solve these problems already!

# THE PROBLEM

- \* Load a dataset containing data from 442 diabetes patients
- \* for each patient there are 10 features (e.g. age, sex, bmi, etc.)
- \* and 1 outcome (“disease progression after one year”)
- \* We’ll be using linear regression to predict disease progression from the 10 features

# LINEAR REGRESSION LAB

- \* If you want to following along, pull the latest version of the **ndap-fa2020** repository from github
- \* <https://github.com/alexhuth/ndap-fa2020/>
- \* Then see **35-linear\_regression-6/35-regression-demos.ipynb**

**END**