

LINEAR REGRESSION VII

11.18.2020

HOMEWORK 5

*** due FRIDAY!!!**

RECAP

- * `np.linalg.lstsq` – numpy function that does least squares regression (often bad)
- * R^2 measures how good a regression model is
- * in-set vs. out-of-set model evaluation
- * regularized (ridge) regression

RECAP: THE PROBLEM


- * Load a dataset containing data from 442 diabetes patients
- * for each patient there are 10 features (e.g. age, sex, bmi, etc.)
- * and 1 outcome (“disease progression after one year”)
- * We’ll be using linear regression to predict disease progression from the 10 features

LINEAR REGRESSION LAB


- * If you want to following along, pull the latest version of the **ndap-fa2020** repository from github
- * <https://github.com/alexhuth/ndap-fa2020/>
- * Then see **35-linear_regression-6/35-regression-demos.ipynb**
- * Today we're picking up where we left off on Monday (see **36-linear_regression-7**)

***WHY IS RIDGE
REGRESSION CHANGING
THE PATTERN OF
WEIGHTS INSTEAD OF
JUST MAKING THE
WEIGHTS SMALLER?***

1D EXAMPLE




$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$



1D EXAMPLE

L2 Regularization:
(as penalty)


$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$



1D EXAMPLE

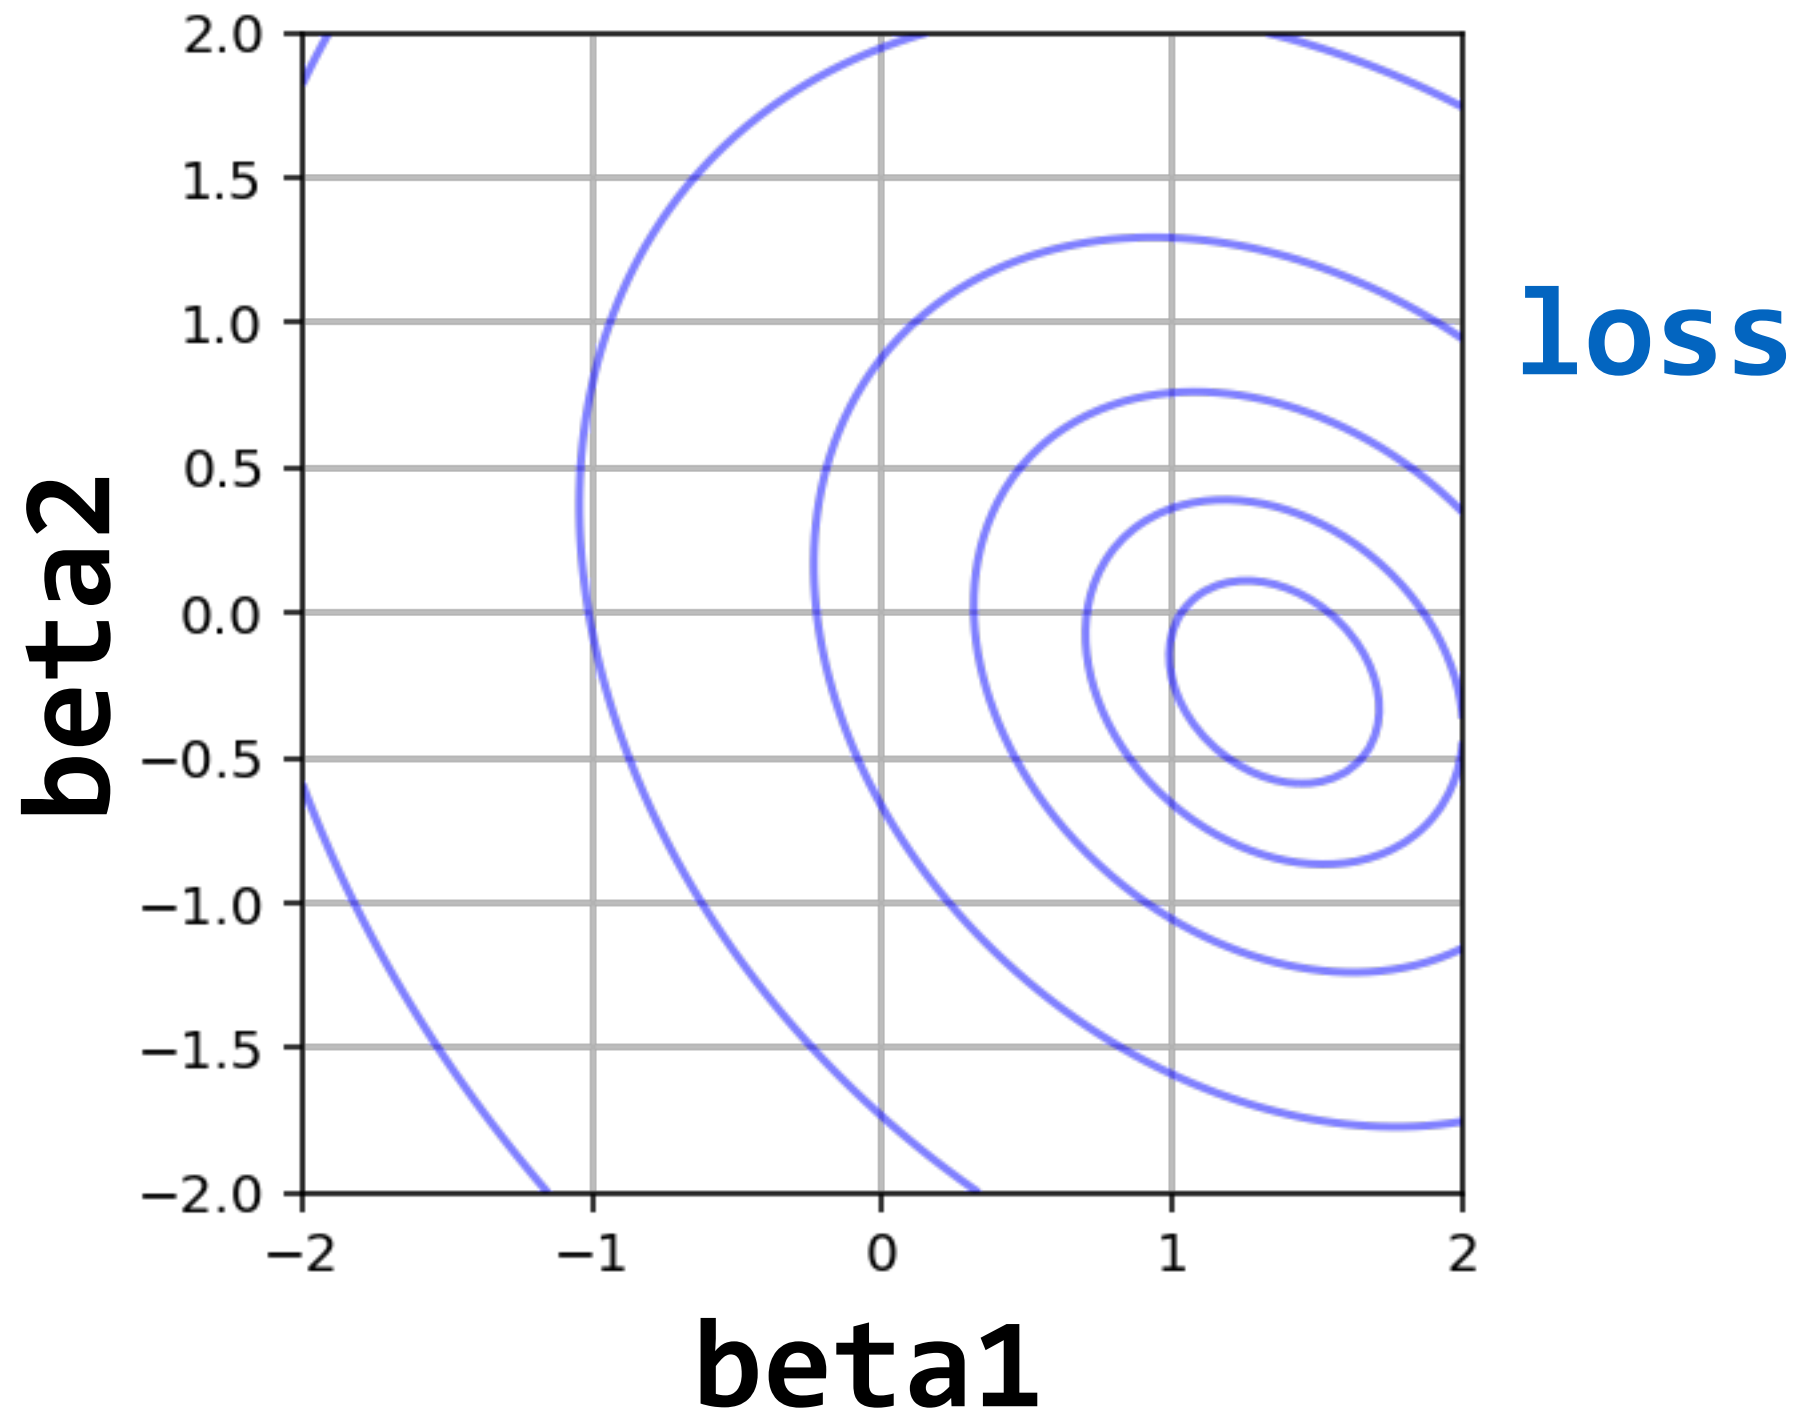
L2 Regularization:
(as penalty)

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$




2D EXAMPLE

$$y_t = x_{1,t}\beta_1 + x_{2,t}\beta_2$$




2D EXAMPLE

$$y_t = x_{1,t}\beta_1 + x_{2,t}\beta_2$$



2D EXAMPLE

$$y_t = x_{1,t}\beta_1 + x_{2,t}\beta_2$$



REGRESSION TAKEAWAYS

- * You want to predict one thing (y) from other things (x_1 , x_2 , etc.)? **Regression!**
- * Regression works best if it's **regularized**
- * You've seen **ridge regression**, but there are many others!
- * Tikhonov regression, sparse regression (LASSO), elastic net regression, etc.

END