


CLASSIFICATION

11.20.2020


PROBLEM SETS

- * **Problem set 5** was due today
- * **Problem set 6** is available today, will be due on December 4



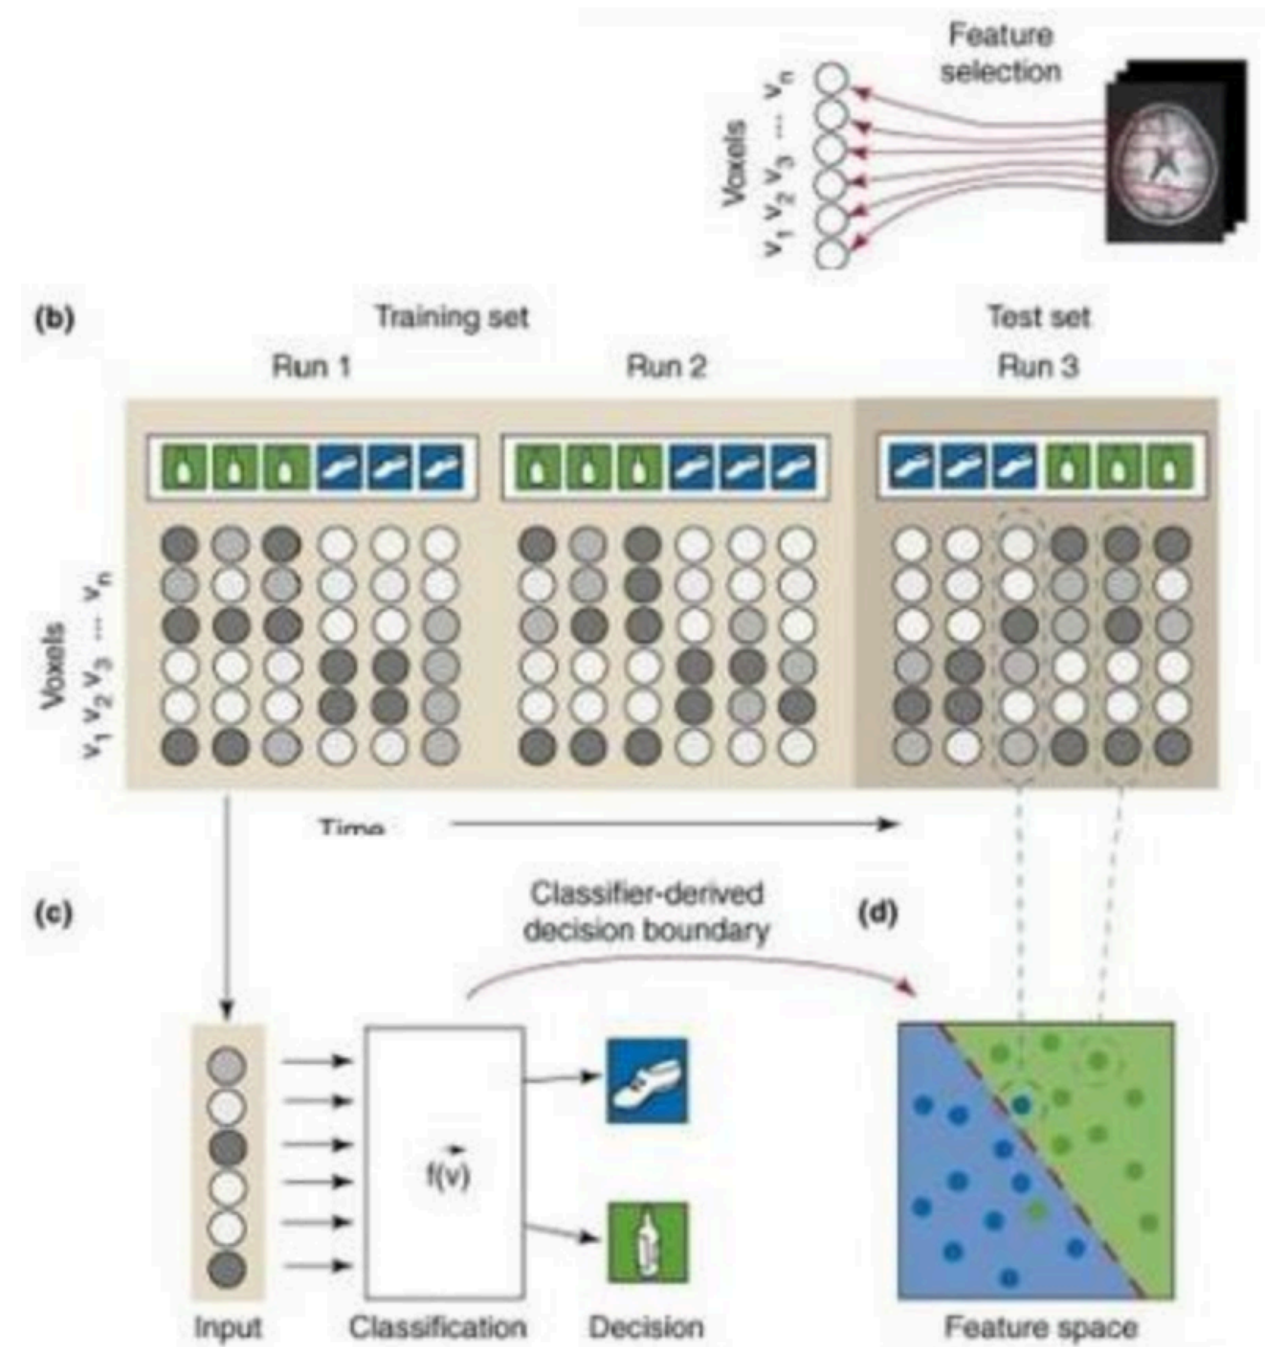
BRAIN-MACHINE INTERFACE

- * suppose you want to build a machine that can help people with **locked-in syndrome**
- * you record EEG activity, and want to classify whether the person is thinking **up** or **down**



FMRI DECODING

- * suppose that you've done an fMRI experiment where people viewed images of SHOES and BOTTLES
- * now you want to use the pattern of fMRI responses to guess whether the stimulus was SHOES or BOTTLES




FMRI DECODING

- * or! your subjects watch videos, and then you guess which objects or actions were present in each video clip based on the responses
- * you want to *classify* each fMRI response into, e.g., “yes there was a face” or “no there was not a face”



CLASSIFIERS





- * a classifier is a function that guesses the “class” of a datapoint based on its features
- * here “class” is **red** or **blue** and each point has 2 features (x and y axes)

CLASSIFIERS

- * there are *MANY* different types of classifiers
- * but they can broadly be divided into two families:
 - * linear classifiers
 - * non-linear classifiers

CLASSIFIERS

- * **linear classifiers (A)** separate the classes by using a straight line (or plane, or hyperplane) as the “decision boundary”
- * **nonlinear classifiers (B)** can separate the classes using an arbitrarily-shaped decision boundary



CLASSIFIERS


- * linear classifiers (e.g. logistic regression, linear SVM) are more appropriate when:
 - * you don't have a lot of data
 - * your data are high-dimensional (many features)
 - * your data are very noisy

CLASSIFIERS

- * nonlinear classifiers (e.g. RBF SVM) are more appropriate when:
 - * your data are low-dimensional (few features)
 - * your data are not very noisy
 - * linear classifiers work poorly :)


TRAINING CLASSIFIERS

- * just like regression models, classifiers should be trained on one dataset, and then tested on an *entirely separate* dataset
- * this controls for *overfitting*



TRAINING CLASSIFIERS

- * like regression models, many classifiers can also be *regularized* (see black curve, below)
- * this can directly reduce overfitting




EVALUATING CLASSIFIERS

- * classifiers can be evaluated by a few metrics
- * one is just the **percent correct**, but that doesn't always tell the whole story
- * a more sensitive way to test classifier performance is using **receiver-operating characteristic (ROC) analysis**


ROC ANALYSIS

- * suppose that your classifier outputs a real number for each datapoint
- * (hopefully) your two classes tend to have different values





ROC ANALYSIS

- * you turn this real value into classification by setting a “criterion value”
- * everything above “criterion” is class A, everything below is class B





ROC ANALYSIS

- * in ROC analysis, you check each possible criterion value
- * and compute the true positive (TP) and false positive (FP) rate for each




ROC ANALYSIS

- * you then plot the TP vs. FP rate to get an ROC curve




ROC ANALYSIS

- * if the classifier is really bad, the ROC should look like a diagonal line
- * if it's really good, the ROC should look like a square



ROC ANALYSIS

- * the area under the ROC curve (AUC) is a common metric
- * AUC tells you how likely the classifier is to correctly order two random points



NEXT TIME

- * on Monday we'll go through some examples of using classification on fMRI data!

THE END