


# CLASSIFICATION

11.20.2020


# PROBLEM SETS

- \* **Problem set 5** was due today
- \* **Problem set 6** is available today, will be due on December 4




# SPIKE SORTING

- \* suppose that you've recorded thousands of spikes from neurons using a tetrode
- \* now you need to figure out which neuron each spike came from based on the recordings from the 4 electrodes




# SPIKE SORTING

- \* you can start by hand-tagging some of the spikes as Neuron A or Neuron B
- \* then you want to automatically *classify* the rest into one of your existing categories



# FMRI DECODING

- \* suppose that you've done an fMRI experiment where people viewed images of SHOES and BOTTLES
- \* now you want to use the pattern of fMRI responses to guess whether the stimulus was SHOES or BOTTLES






# FMRI DECODING

- \* or! your subjects watch videos, and then you guess which objects or actions were present in each video clip based on the responses
- \* you want to *classify* each fMRI response into, e.g., “yes there was a face” or “no there was not a face”



# CLASSIFIERS



- \* a classifier is a function that guesses the “class” of a datapoint based on its features
- \* here “class” is **red** or **blue** and each point has 2 features (x and y axes)



# CLASSIFIERS

- \* there are *MANY* different types of classifiers
- \* but they can broadly be divided into two families:
  - \* linear classifiers
  - \* non-linear classifiers



# CLASSIFIERS

- \* **linear classifiers (A)** separate the classes by using a straight line (or plane, or hyperplane) as the “decision boundary”
- \* **nonlinear classifiers (B)** can separate the classes using an arbitrarily-shaped decision boundary



# CLASSIFIERS


- \* linear classifiers (e.g. logistic regression, linear SVM) are more appropriate when:
  - \* you don't have a lot of data
  - \* your data are high-dimensional (many features)
  - \* your data are very noisy

# CLASSIFIERS

- \* nonlinear classifiers (e.g. RBF SVM) are more appropriate when:
  - \* your data are low-dimensional (few features)
  - \* your data are not very noisy
  - \* linear classifiers work poorly :)


# TRAINING CLASSIFIERS

- \* just like regression models, classifiers should be trained on one dataset, and then tested on an *entirely separate* dataset
- \* this controls for *overfitting*



# TRAINING CLASSIFIERS

- \* like regression models, many classifiers can also be *regularized* (see black curve, below)
- \* this can directly reduce overfitting




# EVALUATING CLASSIFIERS

- \* classifiers can be evaluated by a few metrics
- \* one is just the **percent correct**, but that doesn't always tell the whole story
- \* a more sensitive way to test classifier performance is using **receiver-operating characteristic (ROC) analysis**




# ROC ANALYSIS

- \* suppose that your classifier outputs a real number for each datapoint
- \* (hopefully) your two classes tend to have different values





# ROC ANALYSIS

- \* you turn this real value into classification by setting a “criterion value”
- \* everything above “criterion” is class A, everything below is class B





# ROC ANALYSIS

- \* in ROC analysis, you check each possible criterion value
- \* and compute the true positive (TP) and false positive (FP) rate for each




# ROC ANALYSIS

- \* you then plot the TP vs. FP rate to get an ROC curve




# ROC ANALYSIS

- \* if the classifier is really bad, the ROC should look like a diagonal line
- \* if it's really good, the ROC should look like a square



# ROC ANALYSIS

- \* the area under the ROC curve (AUC) is a common metric
- \* AUC tells you how likely the classifier is to correctly order two random points





# NEXT TIME

- \* on Monday we'll go through some examples of using classification on fMRI data!

**THE END**